

W16_LSTM_MultipleHouses_production

January 6, 2021

1 settings

```
[1]: #settings:  
show_every = 10  
train_for = 100  
learningrate = 1e-3  
stationary = True  
window_size = 168  
houses = [28,37,40,42,105,115,56,51,58,70,99,100]  
reset_scheduler_after_n_epochs = 10
```

2 Initialization

```
[2]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
  
from tqdm import tqdm  
  
from IPython.display import display, HTML  
import time
```

```
[3]: import random  
#Neural Network imports  
import torch  
import torch.nn as nn  
import torch.optim as optim  
  
from sklearn.preprocessing import StandardScaler  
from sklearn.metrics import r2_score  
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import mean_squared_error as mese  
from sklearn.metrics import mean_absolute_error  
  
from torch.utils.data import DataLoader, TensorDataset
```

```
[4]: #cuda imports
ngpu = torch.cuda.device_count() # number of available gpus
device = torch.device("cuda:0") if (torch.cuda.is_available() and ngpu > 0) else "cpu" #cuda:0 for gpu 0, cuda:4 for gpu 5
# torch.backends.cudnn.benchmark=True # Uses cudnn auto-tuner to find the best algorithm to use for your hardware

#Random Seed
random.seed(1337)
torch.manual_seed(1337)

#Scaler objects
scaler_X = StandardScaler()
scaler_y = StandardScaler()
```

Make all functions:

```
[5]: class lstm(nn.Module):
    def __init__(self, feature_size=2, hidden_state_size = 100):
        super().__init__()
        self.hidden_state_size = hidden_state_size
        self.lstm1 = nn.LSTM(feature_size, self.hidden_state_size, batch_first=True)
        self.linear2 = nn.Linear(self.hidden_state_size, 1)

    def forward(self, X): #tensor X
        h, _ = self.lstm1( X ) # h shaped (batch, sequence, hidden_layer)
        h = h[:, -1, :] # only need the output for the last sequence

        y = self.linear2(h) # make a prediction
        y = y + X[:, -1, -1:] # make the output stationary
        return y.view(-1) # like always

    def init_lstm():
        model = lstm().to(device)
        return model
model = init_lstm()
```

```
[6]: #GENERAL FUNCTIONS:
def det(tensor):
    """
    Zet de tensor om van een tensor naar numpy op de CPU.
    """
    return tensor.cpu().detach().numpy()
```

```

def calculate_metrics_for_model(output,target):
    """
    Calculates all the desired evaluation metrics for the model.
    """
    yhat = scaler_y.inverse_transform(det(output))
    y = scaler_y.inverse_transform(det(target))
    actual, pred = np.array(y), np.array(yhat)

    mae = mean_absolute_error(yhat, y)
    mse = mese(yhat, y)
    mape = np.mean(np.abs((actual - pred) / actual)) * 100
    r2 = r2_score(yhat, y)
    return [mae, mse, mape, r2]

def dim3(dft, window=7, gap=24):
    dft = pd.DataFrame(dft)
    #Get time shifted values and apply a moving window
    X = np.concatenate([dft[i:i+window].to_numpy().reshape(1, window, dft.
    ↪shape[1]) for i in range(len(dft)-window-gap)], axis=0)

    #Get the target value (which is the next one in the sequence)
    y = dft.to_numpy()[window + gap:, 1]
    return X.astype(np.float32), y.astype(np.float32)

#LSTM specific functions:
def load_LSTM_data(house_nr):
    """
    Loads the Data for the lstm.
    and returns this.
    """
    house_nr = str(house_nr)
    if len(house_nr)==1:
        house_nr = "00"+str(house_nr)
    if len(house_nr)==2:
        house_nr = "0"+str(house_nr)
    df = pd.read_pickle('/home/18005152/notebooks/zero/Data:/testDataFrames/
    ↪TEST/DeepLearning_production_'+str(house_nr))
    return df

def LSTM_split_df(df):
    """
    Splits the dataframe in train test validate parts.
    """
    trdf = df.loc['2019-01':'2019-08']
    vadf = df.loc['2019-09':'2019-11']
    tedf = df.loc['2019-12':]

```

```

    return trdf, vadf, tedf

def LSTM_data_scaler(df1,df2,df3):
    dftr = df1.copy()
    dfva = df2.copy()
    dfte = df3.copy()

    scaler_X.fit(dftr.iloc[:, :-1])
    scaler_y.fit(dftr.iloc[:, -1:])

    #train
    dftr.iloc[:, :-1] = scaler_X.transform(dftr.iloc[:, :-1])
    dftr.iloc[:, -1:] = scaler_y.transform(dftr.iloc[:, -1:])

    #Valid
    dfva.iloc[:, :-1] = scaler_X.transform(dfva.iloc[:, :-1])
    dfva.iloc[:, -1:] = scaler_y.transform(dfva.iloc[:, -1:])

    #Test
    dfte.iloc[:, :-1] = scaler_X.transform(dfte.iloc[:, :-1])
    dfte.iloc[:, -1:] = scaler_y.transform(dfte.iloc[:, -1:])

    df1 = dftr
    df2 = dfva
    df3 = dfte

    return df1,df2,df3

def LSTM_create_tensors(n1,n2,n3,n4,n5,n6):
    train_X_t = torch.from_numpy(np.array(n1))
    train_y_t = torch.from_numpy(np.array(n2))

    valid_X_t = torch.from_numpy(np.array(n3)).to(device)
    valid_y_t = torch.from_numpy(np.array(n4)).to(device)

    test_X_t = torch.from_numpy(np.array(n5)).to(device)
    test_y_t = torch.from_numpy(np.array(n6)).to(device)
    return train_X_t, train_y_t, valid_X_t, valid_y_t, test_X_t, test_y_t

def train_LSTM_Nepoch(data, N):
    """
    Train the LSTM for one epoch.
    """
    model = init_lstm()
    model.train()
    length_loader = len(data)

```

```

scheduler = torch.optim.lr_scheduler.OneCycleLR(optimizer, u
↪max_lr=learningrate, steps_per_epoch=length_loader, epochs=train_for)
for i in range(N):
    for X, y in data:
        X, y = X.to(device), y.to(device)

        #train LSTM and reshape output:
        optimizer.zero_grad()
        output = model(X)

        #bereken de loss over de output en update de parameters:
        loss = criterion(output, y)
        #save train loss scores correctly (not the last one)

        loss.backward()
        optimizer.step()
        scheduler.step()

    if i % reset_scheduler_after_n_epochs == 0:
        scheduler = torch.optim.lr_scheduler.OneCycleLR(optimizer, u
↪max_lr=learningrate, steps_per_epoch=len(data), u
↪epochs=reset_scheduler_after_n_epochs)
        pass

def validate_LSTM(data, target):
    """
    Validate the LSTM on unseen data and give back the evaluation metrics.
    """
    model.eval()
    optimizer.zero_grad()

    outputV = model(data)
    return calculate_metrics_for_model(outputV,target)

```

3 Main loop:

```

[7]: #stats savelist:
LSTM_stats = pd.DataFrame()

#Training parameters:
optimizer = optim.Adam(model.parameters(), lr=learningrate)
criterion = nn.SmoothL1Loss()

#Learning loop:
for i in tqdm(range(len(houses))):

```

```

house_number = houses[i]
"""
Data loading...
"""

#load the data:
lstm_df = load_LSTM_data(house_number)
#Splits de data in train valid test:
train_LSTM, valid_LSTM, test_LSTM = LSTM_split_df(lstm_df)
#scale de data:
train_LSTM, valid_LSTM, test_LSTM = LSTM_data_scaler(train_LSTM, □
→valid_LSTM, test_LSTM)
#doe een moving window van 3 eroverheen:
train_x_LSTM, train_y_LSTM = dim3(train_LSTM, window_size)
valid_x_LSTM, valid_y_LSTM = dim3(valid_LSTM, window_size)
test_x_LSTM, test_y_LSTM = dim3(test_LSTM, window_size)
#maak tensors van de data:
train_X_t_LSTM, train_y_t_LSTM, valid_X_t_LSTM, valid_y_t_LSTM, □
→test_X_t_LSTM, test_y_t_LSTM = LSTM_create_tensors(train_x_LSTM, □
→train_y_LSTM, valid_x_LSTM, valid_y_LSTM, test_x_LSTM, test_y_LSTM)

"""
Training
"""

train_ds = TensorDataset(train_X_t_LSTM, train_y_t_LSTM)
train_dl = DataLoader(train_ds, batch_size=64, num_workers=3)

target = train_y_t_LSTM.view(-1)
train_LSTM_Nepoch(train_dl,train_for)

model.eval()
output = model(train_X_t_LSTM.to(device))
y = train_y_t_LSTM.to(device)
LSTM_train_stats = calculate_metrics_for_model(output, y)

"""

Evaluation
"""

dataV = valid_X_t_LSTM; targetV = valid_y_t_LSTM.view(-1);
LSTM_valid_stats = validate_LSTM(dataV, targetV)

"""

Testing
"""

dataTest = test_X_t_LSTM; targetTest = test_y_t_LSTM.view(-1);

```

```

LSTM_test_stats = validate_LSTM(dataTest, targetTest)

"""

Save metrics
"""

index = [
    "MAE_train", "MSE_train", "MAPE_train", "R2_train",
    "MAE_valid", "MSE_valid", "MAPE_valid", "R2_valid"
]
New_Stats = pd.DataFrame(LSTM_train_stats + LSTM_valid_stats + LSTM_test_stats,
    index=index, columns=[str(house_number)])
LSTM_stats = pd.concat([LSTM_stats, New_Stats], axis=1)

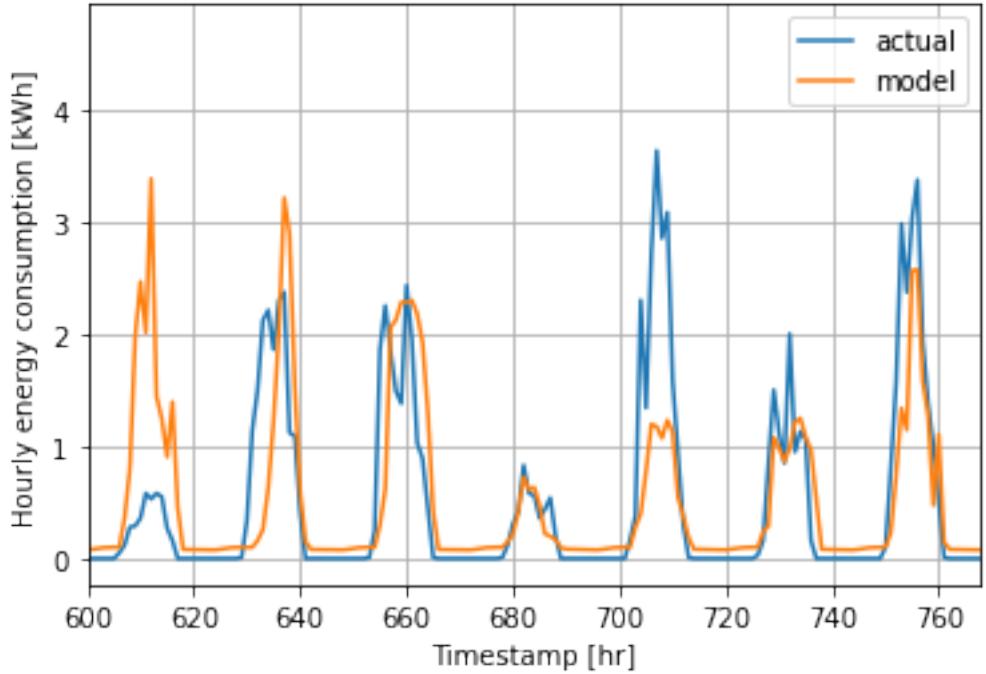
```

100% | 12/12 [26:54<00:00, 134.55s/it]

```
[8]: dataV = valid_X_t_LSTM; targetV = valid_y_t_LSTM.view(-1);
model.eval()
optimizer.zero_grad()

outputV = det(model(dataV))
```

```
[10]: yhat = scaler_y.inverse_transform(outputV)
y = scaler_y.inverse_transform(valid_LSTM.production)
plt.plot(y, label='actual')
plt.plot(yhat, label="model")
plt.xlabel("Timestamp [hr]")
plt.ylabel("Hourly energy consumption [kWh]")
plt.grid()
plt.legend()
plt.xlim([600, 768])
plt.savefig("LSTM_production_house100.png")
```



```
[11]: LSTM_stats
```

	28	37	40	42	\
MAE_train	5.800745e-01	6.290677e-01	5.726565e-01	5.745850e-01	
MSE_train	9.427599e-01	1.094144e+00	9.025427e-01	9.054864e-01	
MAPE_train	1.040593e+08	2.784597e+08	5.490624e+07	1.756825e+09	
R2_train	5.860221e-01	5.800176e-01	5.833813e-01	5.787911e-01	
MAE_valid	3.589388e-01	3.694080e-01	3.414911e-01	3.543031e-01	
MSE_valid	3.887345e-01	4.082521e-01	3.470946e-01	3.719818e-01	
MAPE_valid	1.287012e+08	3.408815e+08	6.718824e+07	2.151421e+09	
R2_valid	4.553477e-01	4.728332e-01	4.712861e-01	4.489117e-01	
MAE_test	1.806626e-01	1.910968e-01	1.675688e-01	1.787679e-01	
MSE_test	7.634789e-02	8.388042e-02	6.237944e-02	7.510534e-02	
MAPE_test	1.490067e+08	4.003416e+08	7.610898e+07	2.447982e+09	
R2_test	1.952975e-01	2.002880e-01	2.281147e-01	1.779697e-01	
	105	115	56	51	\
MAE_train	5.315632e-01	5.793098e-01	5.656943e-01	5.859007e-01	
MSE_train	7.915370e-01	9.202571e-01	8.855981e-01	9.454227e-01	
MAPE_train	9.703863e+07	1.446417e+08	7.156817e+07	1.009758e+08	
R2_train	6.007932e-01	5.817719e-01	5.903437e-01	5.843050e-01	
MAE_valid	3.201535e-01	3.558024e-01	3.325159e-01	3.621002e-01	
MSE_valid	3.059551e-01	3.794333e-01	3.242446e-01	3.911622e-01	
MAPE_valid	1.170376e+08	1.675533e+08	8.487794e+07	1.258250e+08	
R2_valid	4.701279e-01	4.520816e-01	4.724781e-01	4.552206e-01	

MAE_test	1.587741e-01	1.779035e-01	1.728046e-01	1.831995e-01
MSE_test	5.635858e-02	7.474770e-02	6.819513e-02	7.876436e-02
MAPE_test	1.346469e+08	1.952281e+08	9.877236e+07	1.414088e+08
R2_test	2.065491e-01	1.940704e-01	2.065042e-01	1.927188e-01
	58	70	99	100
MAE_train	5.466580e-01	5.831779e-01	5.847489e-01	5.854917e-01
MSE_train	8.333455e-01	9.339937e-01	9.374416e-01	9.424492e-01
MAPE_train	3.607152e+08	1.881661e+08	9.453554e+07	3.974338e+09
R2_train	5.902555e-01	5.776920e-01	5.803581e-01	5.824304e-01
MAE_valid	3.233155e-01	3.625794e-01	3.608990e-01	3.578877e-01
MSE_valid	3.088230e-01	3.902918e-01	3.910239e-01	3.848440e-01
MAPE_valid	4.458169e+08	2.340034e+08	1.137805e+08	4.804672e+09
R2_valid	4.578550e-01	4.475778e-01	4.487678e-01	4.590984e-01
MAE_test	1.692786e-01	1.828210e-01	1.795999e-01	1.725933e-01
MSE_test	6.740455e-02	8.021878e-02	7.518524e-02	6.632595e-02
MAPE_test	5.065594e+08	2.724109e+08	1.289826e+08	5.342758e+09
R2_test	2.012255e-01	1.595703e-01	1.823894e-01	2.083834e-01

[12]: `LSTM_stats.to_pickle("LSTM_statistics")`

[13]: `pd.read_pickle("LSTM_statistics").to_excel("LSTM.xlsx")`

4 summarize stats with mean

[14]: `(LSTM_stats).mean(axis=1)`

```

[14]: MAE_train      5.765774e-01
      MSE_train     9.195815e-01
      MAPE_train    6.021857e+08
      R2_train      5.846801e-01
      MAE_valid     3.499495e-01
      MSE_valid     3.659867e-01
      MAPE_valid    7.318132e+08
      R2_valid      4.592988e-01
      MAE_test       1.762559e-01
      MSE_test      7.207611e-02
      MAPE_test     8.245171e+08
      R2_test       1.960901e-01
      dtype: float64

```